

Finite-Time Convergence of Competitive Graph Recoloring via Discrete Lyapunov Functions

Ranjan Veerabhadraswamy^a, Ajith Jubilson Emerson^{a,*}

^a*School of Computer Science and Engineering, Vellore Institute of Technology (VIT-AP), Andhra Pradesh, India*

Abstract

We study an asynchronous competitive recoloring process on a finite graph, in which each active vertex selfishly changes color only when doing so strictly reduces the number of same-colored neighbors. Although globally minimizing monochromatic edges is equivalent to a MAX- k -CUT objective and is NP-hard in general, the local dynamics admit a simple exact potential. We prove that every accepted recoloring decreases the number of monochromatic edges by exactly the local improvement, and therefore the process reaches a pure Nash equilibrium in at most $|E|$ accepted recolorings. Simulations on sparse Erdős–Rényi and Barabási–Albert graphs support the deterministic bound and show substantially faster average-case convergence.

Keywords: graph coloring, best-response dynamics, potential game, discrete Lyapunov function, MAX- k -CUT

1. Introduction

Graph recoloring under local incentives appears in distributed coordination, interference avoidance, and networked resource assignment [6, 7]. In such settings, a natural objective is to reduce monochromatic edges, but the corresponding global optimization problem is computationally difficult: minimizing monochromatic edges is equivalent to maximizing bichromatic edges, i.e., to the MAX- k -CUT objective, which is NP-hard in general [4, 5]. This

*Corresponding author.

E-mail addresses: ranjanvswamyjnv2005@gmail.com (R. Veerabhadraswamy), ajith.jubilson@vitap.ac.in (A.J. Emerson)

motivates the analysis of distributed local heuristics whose behavior can be certified without solving the global problem.

The game-theoretic view is equally natural. Each vertex is a player, each color is an action, and each player seeks to reduce its local conflict cost. Pure Nash equilibria and potential methods are classical tools for such local incentive systems [1, 2, 3]. The specific rule studied here is a strict best-response recoloring dynamics. The main observation is that the number of monochromatic edges is an exact discrete Lyapunov function, in the spirit of potential-based convergence arguments for networked dynamics [8, 9]. Consequently, selfish local updates cannot cycle and must reach a locally stable coloring after at most $|E|$ accepted moves.

2. Model

Let $G = (V, E)$ be a finite undirected graph, with $n = |V|$ vertices and $m = |E|$ edges. Let $C = \{1, 2, \dots, k\}$ be a finite color set. A configuration is denoted by

$$c = (c_i)_{i \in V} \in C^V,$$

where c_i is the color assigned to vertex i . For $i \in V$, let $N(i) = \{j \in V : \{i, j\} \in E\}$ be its neighborhood. The local conflict cost of vertex i is

$$d_i(c_i, c_{-i}) = \sum_{j \in N(i)} \mathbf{1}\{c_i = c_j\},$$

where c_{-i} denotes the colors of all vertices except i . Equivalently, if the colors of the other vertices are fixed, then for any $\gamma \in C$,

$$d_i(\gamma, c_{-i}) = \sum_{j \in N(i)} \mathbf{1}\{\gamma = c_j\}.$$

The dynamics are asynchronous. At each activation, a single vertex i is selected. If there exists a color that strictly reduces its local conflict cost, vertex i changes to a best-response color

$$c_i^+ \in \arg \min_{\gamma \in C} d_i(\gamma, c_{-i}),$$

and the update is accepted only when

$$d_i(c_i^+, c_{-i}) < d_i(c_i, c_{-i}).$$

Otherwise the configuration is unchanged. In what follows, a *step* means one accepted recoloring.

Define the global conflict potential

$$\Phi(c) = \sum_{\{u,v\} \in E} \mathbf{1}\{c_u = c_v\} = \frac{1}{2} \sum_{i \in V} d_i(c_i, c_{-i}).$$

Thus $\Phi(c)$ is exactly the number of monochromatic edges.

3. Finite convergence

Lemma 1 (Exact potential drop). *Let c and c^+ be two configurations that differ only at one vertex i . Suppose vertex i changes from color $\alpha = c_i$ to color $\beta = c_i^+$ and strictly reduces its local conflict cost, while $c_{-i}^+ = c_{-i}$. Then*

$$\Phi(c^+) - \Phi(c) = d_i(\beta, c_{-i}) - d_i(\alpha, c_{-i}) \leq -1.$$

Proof. Only edges incident to i can change their conflict status, since all other vertices retain their colors. Hence every edge not in $\{\{i, j\} : j \in N(i)\}$ contributes the same amount to $\Phi(c^+)$ and $\Phi(c)$.

Before the update, the number of conflicting edges incident to i is

$$d_i(\alpha, c_{-i}) = \sum_{j \in N(i)} \mathbf{1}\{\alpha = c_j\}.$$

After the update, the number of conflicting edges incident to i is

$$d_i(\beta, c_{-i}) = \sum_{j \in N(i)} \mathbf{1}\{\beta = c_j\}.$$

Therefore

$$\Phi(c^+) - \Phi(c) = d_i(\beta, c_{-i}) - d_i(\alpha, c_{-i}).$$

The strict-improvement rule gives $d_i(\beta, c_{-i}) < d_i(\alpha, c_{-i})$. Since both quantities are integers, their difference is at most -1 . $\blacksquare$

Theorem 1 (Finite convergence). *For any finite graph $G = (V, E)$, any finite color set C , and any initial configuration $c(0) \in C^V$, the competitive recoloring dynamics terminate after at most $|E|$ accepted recoloring steps. More precisely, if T denotes the number of accepted recolorings before termination, then*

$$T \leq \Phi(c(0)) \leq |E|.$$

The terminal configuration is a pure Nash equilibrium, equivalently a fixed point of the strict best-response dynamics.

Proof. By Lemma 1, every accepted recoloring decreases Φ by at least one. The potential is integer-valued and nonnegative because it counts monochromatic edges:

$$0 \leq \Phi(c) \leq |E| \quad \text{for every } c \in C^V.$$

Thus a sequence of accepted recolorings can have length no larger than the initial potential value. Consequently,

$$T \leq \Phi(c(0)) \leq |E|.$$

No infinite sequence of accepted recolorings is possible.

At termination, no vertex has a color in C that strictly lowers its local conflict cost. Hence no individual vertex can improve by a unilateral recoloring while the other vertices remain fixed. This is precisely a pure Nash equilibrium for the local conflict-cost game and a fixed point of the stated dynamics. $\blacksquare$

4. Empirical validation

Theorem 1 gives a deterministic worst-case upper bound. To illustrate its practical scale, we simulated randomized asynchronous updates on Erdős–Rényi and Barabási–Albert graphs [10, 11] with $k = 3$ colors and constant average degree 8. For each family and each $n \in \{100, 200, 400, 800\}$, we ran 1000 independent trials, giving 8000 runs in total.

No run failed to converge, and no run violated the bound $T \leq \Phi(c(0)) \leq |E|$. The largest observed value of $T/|E|$ was 0.1779, while the average across graph sizes and families was approximately 0.121. Since these graph families are sparse, $|E| = \mathcal{O}(|V|)$; the observed convergence is therefore essentially linear in the network size for these experiments.

Remark 1. *The observed speedup is consistent with the randomized asynchronous schedule: activation orders that approach the deterministic worst case are rarely sampled in standard sparse graph ensembles. The experiments suggest behavior closer to $\mathcal{O}(|V| \log |V|)$, and nearly linear in the tested range, but the formal guarantee used here remains the deterministic edge bound in Theorem 1.*




Figure 1: Accepted recoloring steps versus the deterministic edge bound. Each point is one trial on a sparse Erdős-Rényi or Barabási-Albert graph. The dashed line is $T = |E|$. The empirical trajectories remain far below the worst-case Lyapunov bound.

5. Conclusion

Competitive graph recoloring admits an exact discrete Lyapunov function: the number of monochromatic edges. This yields a concise finite-time convergence proof and a deterministic accepted-move bound of $T \leq \Phi(c(0)) \leq |E|$. The simulations confirm the bound and indicate much faster typical convergence on standard sparse random graph families.

The result is intentionally local: it does not require a central coordinator, global optimization oracle, or assumptions on the activation order beyond single-vertex asynchronous updates. Since global minimization of conflicts is intractable in general, the value of the theorem is to certify that a purely selfish rule still has a finite, monotone path to equilibrium. This separates the achievable distributed guarantee from the harder question of finding a globally optimal coloring, and it explains why the empirical behavior can be studied cleanly through the same potential.

Acknowledgements. The authors thank their mentors, colleagues, and peers for encouragement and helpful feedback. This research received no external funding. The simulation code is available from the authors upon reasonable request.

References

- [1] J. Nash, Non-cooperative games, *Annals of Mathematics* 54 (1951) 286–295.
- [2] R. W. Rosenthal, A class of games possessing pure-strategy Nash equilibria, *International Journal of Game Theory* 2 (1973) 65–67.
- [3] D. Monderer, L. S. Shapley, Potential games, *Games and Economic Behavior* 14 (1996) 124–143.
- [4] M. R. Garey, D. S. Johnson, *Computers and Intractability: A Guide to the Theory of NP-Completeness*, W. H. Freeman, 1979.
- [5] A. Frieze, M. Jerrum, Improved approximation algorithms for MAX k -CUT and MAX BISECTION, *Algorithmica* 18 (1997) 67–81.
- [6] F. Kuhn, R. Wattenhofer, On the complexity of distributed graph coloring, in: *Proceedings of the 25th Annual ACM Symposium on Principles of Distributed Computing*, 2006, pp. 7–15.
- [7] P. N. Panagopoulou, P. G. Spirakis, A game theoretic approach for efficient graph coloring, in: *Proceedings of the 19th International Symposium on Algorithms and Computation*, 2008, pp. 183–195.
- [8] J. J. Hopfield, Neural networks and physical systems with emergent collective computational abilities, *Proceedings of the National Academy of Sciences* 79 (1982) 2554–2558.
- [9] K. R. Apt, S. Simon, A classification of weakly acyclic games, in: *Proceedings of the 5th International Symposium on Algorithmic Game Theory*, 2012, pp. 1–12.
- [10] P. Erdős, A. Rényi, On random graphs. I, *Publicationes Mathematicae* 6 (1959) 290–297.
- [11] A.-L. Barabási, R. Albert, Emergence of scaling in random networks, *Science* 286 (1999) 509–512.